

Venkata Sai Sushma Emmadi

sushmaevs@gmail.com | +1 (667) 445-8558 | [linkedin.com/in/sushma-evs](https://www.linkedin.com/in/sushma-evs) | <https://github.com/sushma62216>

PROFESSIONAL SUMMARY

Data Analyst with experience translating product and operational data into dashboards, KPI frameworks, and insights for cross-functional teams. Strong in SQL and Python, with working knowledge of Power BI, Tableau, and cloud platforms (AWS, Azure). Partnering with stakeholders to define measurement strategy, validate data quality, and communicate findings clearly across technical and non-technical audiences

TECHNICAL SKILLS

Analytics & Visualization: Python (Pandas, NumPy, Matplotlib), SQL, Tableau, Power BI

Data Analysis: Statistical Analysis, KPI Reporting, Ad Hoc Analysis, Data Cleaning, Data Wrangling

Business Intelligence: Dashboard Development, Data Storytelling, Data Visualization, Reporting

Data Engineering: Apache Airflow, dbt, Spark, Docker, ETL Pipelines

Databases & Cloud: Snowflake, PostgreSQL, MySQL, Microsoft Azure

Version Control: Git, GitHub

Certifications: Microsoft Certified: Azure Fundamentals (2025)

PROFESSIONAL EXPERIENCE

Data Analyst

April 2026 – Present

Kahana | Chicago, Illinois (Remote)

- Defining the product analytics framework for the upcoming Product Hunt launch of Oasis (AI-powered productivity tool), scoping 20+ KPIs across activation, retention, monetization, and feedback quality to establish measurement baselines for an early-stage startup.
- Mapping the company's Supabase data model to planned analytics use cases — identifying which events, tables, and user actions need to be tracked to support user behavior analysis, conversion funnels, and product usage monitoring post-launch.
- Collaborating with the founding team to translate product goals into measurable success criteria, including feedback quality monitoring and anomaly detection logic to flag low-signal user submissions in training data.

Data Analyst

May 2022 – May 2023

Appleute | Hammelburg, Germany

- Built data pipelines to ingest and transform 100K+ records, implementing data validation and quality checks to ensure reliable, analytics-ready datasets for downstream reporting and decision-making
- Developed an anomaly detection system applying statistical methods to validate 50K+ data points, and built data visualizations in Python that reduced ad hoc analysis time by 30%, improving data trust for business stakeholders
- Collaborated with cross-functional teams to translate business questions into structured data preprocessing workflows, cleaning and transforming 10K+ records to meet quality standards for downstream analysis

- Optimized SQL queries powering analytics dashboards, improving data retrieval performance by 40% and enabling faster access to key product and business metrics for stakeholders and decision-makers

PROJECTS

Skyra-Atmos: Climate Data Analytics Pipeline

Feb 2026

Independent Project

- Designed an end-to-end data pipeline for ingesting and transforming large-scale datasets (20+ years, 11K+ stations) from AWS S3 into a Snowflake data warehouse, applying dimensional modeling and ER diagrams to structure fact and dimension tables
- Implemented data modeling with RAW/STAGING/ANALYTICS schemas, following best practices for data warehouse architecture, data marts, and documentation
- Orchestrated automated ETL workflows using Apache Airflow DAGs, implementing incremental load logic for Parquet data to keep the warehouse in sync with source systems and reduce full-refresh runtimes
- Built Power BI dashboards on top of the analytics layer to surface station-level trends, patterns, and KPIs, translating warehouse data into clear, decision-ready insights for end users

Medicare Part D 2023 — Prescriber Spend Intelligence

Apr 2026

Independent Project

- Shipped Phase 1 data pipeline ingesting 1.1M CMS Medicare prescriber records into Snowflake and transforming them through a dbt staging to marts layer, producing five aggregate mart tables (specialty, drug, state, top providers, specialty-drug) ready for downstream analytics and reporting
- Designed a pre-aggregated embedding strategy that converts ~10K–50K dbt mart rows into natural-language summary strings — cutting the embedding corpus ~95% vs. row-level embedding while preserving accuracy for aggregate queries
- Modeled the full raw to staging and marts lineage in dbt with schema tests on key columns (NPI, drug name) and idempotent transformations for safe reruns
- Currently building Phase 2 (RAG + app layer): in-database vector embeddings and cosine-similarity retrieval natively on Snowflake, a FastAPI backend powered by Mistral LLM via Snowflake Cortex, and a Next.js chat UI returning answers with cited sources and multi-turn conversation memory

EDUCATION

Master of Professional Studies — Data Science

Jan 2024 – Dec 2025

University of Maryland, Baltimore County (UMBC) | GPA: 3.8 / 4.0

Relevant Coursework: Big Data Platforms, Machine Learning, Statistical Analysis, Data Management